

TMDb

July 8, 2018

1 Project: Exploring the Movie Database (TMDB)

1.1 Table of contents

Introduction

Wrangling

Exploratory data analysis

Conclusions

Introduction

I like action, animation, fantasy, and science fiction movies (not necessarily in that order!). For a personal project I would have liked to know what are the most popular movies in these categories. Do other people like these categories as well? Do these movies make money?

But that would be a personal project. I have refrained from these questions in the current project. Instead I have followed the standard procedure of wrangling and exploring a given dataset. Along the way I have asked questions based upon the features of the dataset.

1. What are the most popular movies of all time?
2. What are the features associated with these movies?
3. Who are the most popular directors?
4. How is the revenue distribution?
5. What are the top-grossing movies of all time?
6. Who are the directors of top-grossing movies?
7. Is popularity related to revenue?
8. Who are the most productive directors?
9. What's the yearly movie production rate?
10. Which genres have been popular over the years?

In the following, I first describe the data wrangling phase where I load the data into a dataframe, and then assess and clean. Then I describe my explorations of the dataset. I have explored *popularity*, *revenue*, *their interrelation*, *directors*, *years*, and* *genres*.* Along the way I also find the answers to the questions posed above.

1.1.1 Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
% matplotlib inline
```

```
import seaborn as sns
sns.set()
```

Wrangle

1.1.2 Gather

```
In [2]: dataset = pd.read_csv('tmdb-movies.csv');
dataset.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	

```

2 Beatrice Prior must confront her inner demons ...      119
3 Thirty years after defeating the Galactic Empi...      136
4 Deckard Shaw seeks revenge against Dominic Tor...      137

                                genres \
0 Action|Adventure|Science Fiction|Thriller
1 Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller
3 Action|Adventure|Science Fiction|Fantasy
4      Action|Crime|Thriller

                                production_companies release_date vote_count \
0 Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1 Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2 Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480
3      Lucasfilm|Truenorth Productions|Bad Robot      12/15/15      5292
4 Universal Pictures|Original Film|Media Rights ...      4/1/15      2947

    vote_average  release_year    budget_adj    revenue_adj
0           6.5         2015  1.379999e+08  1.392446e+09
1           7.1         2015  1.379999e+08  3.481613e+08
2           6.3         2015  1.012000e+08  2.716190e+08
3           7.5         2015  1.839999e+08  1.902723e+09
4           7.3         2015  1.747999e+08  1.385749e+09

[5 rows x 21 columns]

```

We see that there are 21 features associated with each entry. The questions I am interested in primarily concerns the popularity and box-office performance of a movie. In this respect, some of the features like 'id', 'imdb_id', and 'homepage' are irrelevant. In the next phases of the wrangling procedure I shall consider modifying the dataframe to our need.

1.1.3 Assess

Because a number of features could not be displayed abobe, let's first generate a list of the features in the dataset. This information will be helpful in deciding which features to keep.

```
In [3]: dataset.columns
```

```
Out[3]: Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',
              'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',
              'runtime', 'genres', 'production_companies', 'release_date',
              'vote_count', 'vote_average', 'release_year', 'budget_adj',
              'revenue_adj'],
              dtype='object')
```

We can safely drop 'id', 'imdb_id', 'budget', 'revenue', and 'homepage' features: 'id' , 'imdb_id', and 'homepage' should have any relevance to the popularity and revenue of a movie, and raw 'budget' and 'revenue' are unnecessary as their adjusted values are also given in the

final columns. It may be argued that the tagline, keywords, and overview matters for the popularity and revenue of a movie. However, this information may be redundant when the genres are specified or just too difficult to use. In the following, I have therefore decided to drop these features as well. Because the popularity score is already given, I shall not need 'vote_count' and 'vote_average'.

```
In [4]: list_drop = ['id', 'imdb_id', 'budget', 'revenue', 'homepage', 'tagline', 'overview', 'v
dataset.drop(list_drop, axis=1, inplace=True)
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 11 columns):
popularity                10866 non-null float64
original_title            10866 non-null object
cast                      10790 non-null object
director                  10822 non-null object
runtime                   10866 non-null int64
genres                    10843 non-null object
production_companies      9836 non-null object
release_date              10866 non-null object
release_year              10866 non-null int64
budget_adj                10866 non-null float64
revenue_adj               10866 non-null float64
dtypes: float64(3), int64(2), object(6)
memory usage: 933.9+ KB
```

1.1.4 Clean

Time to clean data! We see that some entries for cast are missing. Let's have a look at some of these movies.

Cast

```
In [5]: dataset[dataset.cast.isnull()].head()
```

```
Out[5]:
```

	popularity	original_title	cast	\
371	0.422901	Sanjay's Super Team	NaN	
441	0.220751	Winter on Fire: Ukraine's Fight for Freedom	NaN	
465	0.201696	Bitter Lake	NaN	
536	0.122543	A Faster Horse	NaN	
538	0.114264	The Mask You Live In	NaN	

	director	runtime	genres	\
371	Sanjay Patel	7	Animation	
441	Evgeny Afineevsky	98	Documentary	
465	Adam Curtis	135	Documentary	
536	David Gelb	90	Documentary	

538 Jennifer Siebel Newsom 88 Documentary

	production_companies	release_date	\
371	Pixar Animation Studios	11/25/15	
441	Passion Pictures Campbell Grobman Films Afinee...	10/9/15	
465	BBC	1/24/15	
536	NaN	10/8/15	
538	NaN	1/1/15	

	release_year	budget_adj	revenue_adj
371	2015	0.0	0.0
441	2015	0.0	0.0
465	2015	0.0	0.0
536	2015	0.0	0.0
538	2015	0.0	0.0

```
In [6]: dataset[dataset.cast.isnull()].shape
```

```
Out[6]: (76, 11)
```

We see that adjusted budget and revenue seem to be missing as well. (They cannot be zero!) Moreover, there are just 76 such items out of 10754, less than one percent. I therefore drop these entries in the following.

```
In [7]: dataset = dataset[pd.notnull(dataset['cast'])]
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10790 entries, 0 to 10865
Data columns (total 11 columns):
popularity          10790 non-null float64
original_title      10790 non-null object
cast                10790 non-null object
director            10752 non-null object
runtime             10790 non-null int64
genres              10768 non-null object
production_companies 9800 non-null object
release_date        10790 non-null object
release_year        10790 non-null int64
budget_adj          10790 non-null float64
revenue_adj         10790 non-null float64
dtypes: float64(3), int64(2), object(6)
memory usage: 1011.6+ KB
```

Director The next missing field is 'director'. Let's have a look at the missing entries.

```
In [8]: dataset[dataset.director.isnull()].head()
```

```

Out[8]:      popularity      original_title \
532      0.126594      Iliza Shlesinger: Freezing Hot
548      0.108072      Sense8: Creating the World
556      0.100910      With This Ring
1032      0.291253      Marvel Studios: Assembling a Universe
1054      0.269468      Unlocking Sherlock

      cast director runtime \
532      Iliza Shlesinger      NaN      71
548      Tuppencc Middleton|Bae Doona |Brian J. Smith|A...      NaN      25
556      Regina Hall|Jill Scott|Eve|Brooklyn Sudano|Dei...      NaN      105
1032      Robert Downey Jr.|Chris Hemsworth|Chris Evans|...      NaN      43
1054      Benedict Cumberbatch|Martin Freeman|Steven Mof...      NaN      60

      genres \
532      Comedy
548      Documentary|Science Fiction
556      Comedy|Romance
1032      TV Movie|Documentary
1054      TV Movie|Documentary

      production_companies release_date release_year \
532      New Wave Entertainment      1/23/15      2015
548      Netflix      8/10/15      2015
556      Lifetime Television|Sony Pictures Television      1/24/15      2015
1032      Marvel Studios|ABC Studios      3/18/14      2014
1054      NaN      1/19/14      2014

      budget_adj revenue_adj
532      0.0      0.0
548      0.0      0.0
556      0.0      0.0
1032      0.0      0.0
1054      0.0      0.0

```

```
In [9]: dataset[dataset.director.isnull()].shape
```

```
Out[9]: (38, 11)
```

Budget and revenue are missing for these ones as well. There are just 38 of them. I shall remove them too.

```
In [10]: dataset = dataset[pd.notnull(dataset['director'])]
dataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10752 entries, 0 to 10865
Data columns (total 11 columns):
popularity      10752 non-null float64

```

```

original_title      10752 non-null object
cast                10752 non-null object
director            10752 non-null object
runtime             10752 non-null int64
genres              10732 non-null object
production_companies 9780 non-null object
release_date        10752 non-null object
release_year        10752 non-null int64
budget_adj          10752 non-null float64
revenue_adj         10752 non-null float64
dtypes: float64(3), int64(2), object(6)
memory usage: 1008.0+ KB

```

Genre Some genre entries are missing.

```
In [11]: dataset[dataset.genres.isnull()].head()
```

```

Out[11]:
   popularity  original_title \
424    0.244648    Belli di papà
997    0.330431    Star Wars Rebels: Spark of Rebellion
1712   0.302095    Prayers for Bobby
1897   0.020701    Jonas Brothers: The Concert Experience
2370   0.081892    Freshman Father

   cast \
424    Diego Abatantuono|Matilde Gioli|Andrea Pisani|...
997    Freddie Prinze Jr.|Vanessa Marshall|Steve Blum...
1712   Ryan Kelley|Sigourney Weaver|Henry Czerny|Dan ...
1897   Nick Jonas|Joe Jonas|Kevin Jonas|John Lloyd Ta...
2370   Britt Irvin|Merrilyn Gann|Barbara Tyson|Anthon...

   director  runtime  genres  production_companies \
424    Guido Chiesa    100   NaN                  NaN
997    Steward Lee|Steven G. Lee    44   NaN                  NaN
1712    Russell Mulcahy    88   NaN    Daniel Sladek Entertainment
1897    Bruce Hendricks    76   NaN                  NaN
2370    Michael Scott     0   NaN                  NaN

   release_date  release_year  budget_adj  revenue_adj
424    10/29/15         2015         0.0         0.0
997    10/3/14         2014         0.0         0.0
1712    2/27/09         2009         0.0         0.0
1897    2/27/09         2009         0.0         0.0
2370    6/5/10         2010         0.0         0.0

```

```
In [12]: dataset[dataset.genres.isnull()].shape
```

```
Out[12]: (20, 11)
```

And these too have missing budget and revenue and are few in numbers. Drop them!

```
In [13]: dataset = dataset[pd.notnull(dataset['genres'])]
        dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10732 entries, 0 to 10865
Data columns (total 11 columns):
popularity                10732 non-null float64
original_title            10732 non-null object
cast                     10732 non-null object
director                 10732 non-null object
runtime                  10732 non-null int64
genres                   10732 non-null object
production_companies      9773 non-null object
release_date             10732 non-null object
release_year             10732 non-null int64
budget_adj               10732 non-null float64
revenue_adj              10732 non-null float64
dtypes: float64(3), int64(2), object(6)
memory usage: 1006.1+ KB
```

Production company Curiously a significant number of movies have their production company missing. Let's take a look.

```
In [14]: dataset[dataset.production_companies.isnull()].head()
```

```
Out[14]:
```

	popularity	original_title \	cast	director \	runtime	genres	production_companies	release_date \
228	0.584363	Racing Extinction			90	Adventure Documentary	NaN	1/24/15
259	0.476341	Crown for Christmas		Alex Zamm	84	TV Movie	NaN	11/27/15
295	0.417191	12 Gifts of Christmas		Peter Sullivan	84	Family TV Movie	NaN	11/26/15
298	0.370258	The Girl in the Photographs		Nick Simon	95	Crime Horror Thriller	NaN	9/14/15
328	0.367617	Advantageous		Jennifer Phang	92	Science Fiction Drama Family	NaN	6/23/15

	release_year	budget_adj	revenue_adj
228	2015	0.0	0.0
259	2015	0.0	0.0
295	2015	0.0	0.0
298	2015	0.0	0.0
328	2015	0.0	0.0

```
In [15]: dataset[dataset.production_companies.isnull()].shape
```

```
Out[15]: (959, 11)
```


While the budget and revenue information are missing for these ones as well, they are quite large in number. I shall keep them for now. I shall disregard them while exploring revenue.

Exploratory data analysis

1.1.5 1. Popularity

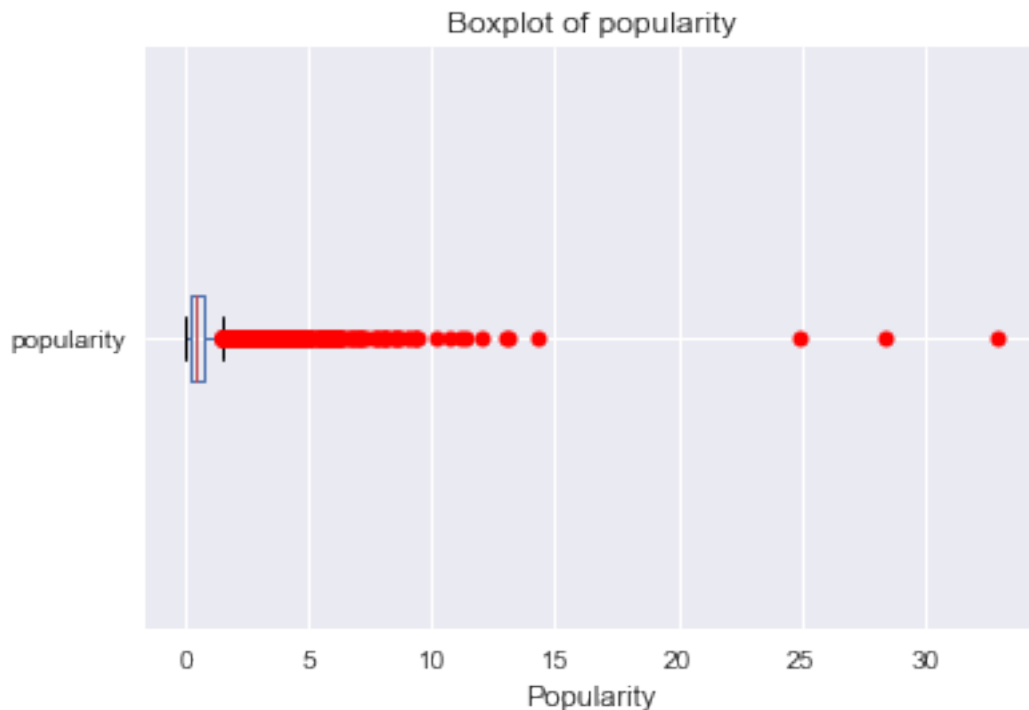
Let's start with popularity. Because it is a float, we start with histogram analysis. Our goal here is to see the distribution of popularity across movies.

```
In [16]: dataset.popularity.hist()
plt.xlabel('Popularity')
plt.ylabel('Counts')
plt.title('Distribution of popularity');
```



It seems like there are a outliers which have skewed the distribution. To confirm this I use boxplot in the following with the "definite" outliers marked red.

```
In [17]: dataset.popularity.plot(kind='box', vert=False, sym='r')
plt.xlabel('Popularity')
plt.title('Boxplot of popularity');
```



This confirms my suspicion about outliers. The outliers are the most popular movies in the database. For the time-being I am interested in the distribution of *typical* popularity values. To identify the typical points in shall use the five-point summary.

```
In [18]: dataset.popularity.describe()
```

```
Out[18]: count      10732.000000
mean         0.652609
std          1.004757
min          0.000188
25%          0.210766
50%          0.387136
75%          0.720621
max          32.985763
Name: popularity, dtype: float64
```

Regular We shall now determine inner fences in boxplot to select the typical (regular) values.


```
In [19]: # Quartiles and fences
q1_pop = dataset.popularity.describe()['25%']
q3_pop = dataset.popularity.describe()['75%']
iqr_pop = q3_pop - q1_pop
fence_low_pop = q1_pop - 1.5 * iqr_pop
fence_high_pop = q3_pop + 1.5 * iqr_pop

# Regular and outlier samples
dataset_pop_reg = dataset.query('popularity > {} and popularity < {}'.format(fence_low_pop, fence_high_pop))
dataset_pop_out = dataset.query('popularity < {} or popularity > {}'.format(fence_low_pop, fence_high_pop))
dataset_pop_reg.shape
```

Out[19]: (9796, 11)

We now plot the histogram of typical values of popularity.

```
In [20]: dataset_pop_reg.popularity.plot.hist(bins=20)
plt.xlabel('Popularity')
plt.title('Distribution of "regular" values of popularity');
```



We see that the typical value of popularity is around 0.3. It is however heavily skewed to the right, meaning that more than half of the movies have popularity greater than this. To get more exact figures we generate the statistics.

```
In [21]: dataset_pop_reg.popularity.describe()
```

```
Out[21]: count      9796.000000
         mean        0.438906
         std         0.325973
         min         0.000188
         25%         0.198252
         50%         0.351092
         75%         0.595587
         max         1.483329
         Name: popularity, dtype: float64
```

Top 10! What are the top ten movies in terms of popularity? This relates to our first few questions. In the following we shall explore features associated with the most popular movies in the list.

```
In [22]: dataset.sort_values(by='popularity', ascending=False).head(10)
```

```
Out[22]:
```

	popularity	original_title \
0	32.985763	Jurassic World
1	28.419936	Mad Max: Fury Road
629	24.949134	Interstellar
630	14.311205	Guardians of the Galaxy
2	13.112507	Insurgent
631	12.971027	Captain America: The Winter Soldier
1329	12.037933	Star Wars
632	11.422751	John Wick
3	11.173104	Star Wars: The Force Awakens
633	10.739009	The Hunger Games: Mockingjay - Part 1

	cast \
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...
629	Matthew McConaughey Jessica Chastain Anne Hath...
630	Chris Pratt Zoe Saldana Dave Bautista Vin Dies...
2	Shailene Woodley Theo James Kate Winslet Ansel...
631	Chris Evans Scarlett Johansson Sebastian Stan ...
1329	Mark Hamill Harrison Ford Carrie Fisher Peter ...
632	Keanu Reeves Michael Nyqvist Alfie Allen Wille...
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...
633	Jennifer Lawrence Josh Hutcherson Liam Hemswor...

	director	runtime \
0	Colin Trevorrow	124
1	George Miller	120
629	Christopher Nolan	169
630	James Gunn	121
2	Robert Schwentke	119
631	Joe Russo Anthony Russo	136
1329	George Lucas	121

632	Chad Stahelski David Leitch	101
3	J.J. Abrams	136
633	Francis Lawrence	123

	genres \
0	Action Adventure Science Fiction Thriller
1	Action Adventure Science Fiction Thriller
629	Adventure Drama Science Fiction
630	Action Science Fiction Adventure
2	Adventure Science Fiction Thriller
631	Action Adventure Science Fiction
1329	Adventure Action Science Fiction
632	Action Thriller
3	Action Adventure Science Fiction Fantasy
633	Science Fiction Adventure Thriller

	production_companies	release_date \
0	Universal Studios Amblin Entertainment Legenda...	6/9/15
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15
629	Paramount Pictures Legendary Pictures Warner B...	11/5/14
630	Marvel Studios Moving Picture Company (MPC) Bu...	7/30/14
2	Summit Entertainment Mandeville Films Red Wago...	3/18/15
631	Marvel Studios	3/20/14
1329	Lucasfilm Twentieth Century Fox Film Corporation	3/20/77
632	Thunder Road Pictures Warner Bros. 87Eleven De...	10/22/14
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15
633	Lionsgate Color Force	11/18/14

	release_year	budget_adj	revenue_adj
0	2015	1.379999e+08	1.392446e+09
1	2015	1.379999e+08	3.481613e+08
629	2014	1.519800e+08	5.726906e+08
630	2014	1.565855e+08	7.122911e+08
2	2015	1.012000e+08	2.716190e+08
631	2014	1.565855e+08	6.583651e+08
1329	1977	3.957559e+07	2.789712e+09
632	2014	1.842182e+07	7.252661e+07
3	2015	1.839999e+08	1.902723e+09
633	2014	1.151364e+08	6.927528e+08

The list definitely makes sense (although I was somewhat surprised by the entry John Wick!). This also answers the first question: **What are the most popular movies of all time?** In the following we explore the numerical features of these movies to answer our second question: **What are the features associated with these movies?**

```
In [23]: dataset.sort_values(by='popularity', ascending=False).head(10).describe()
```

```
Out[23]:
```

	popularity	runtime	release_year	budget_adj	revenue_adj
count	10.000000	10.000000	10.000000	1.000000e+01	1.000000e+01

mean	17.212237	127.000000	2010.700000	1.199485e+08	9.413288e+08
std	8.274145	17.676098	11.851395	5.334061e+07	8.452280e+08
min	10.739009	101.000000	1977.000000	1.842182e+07	7.252661e+07
25%	11.576547	120.250000	2014.000000	1.046841e+08	4.042936e+08
50%	13.041767	122.000000	2014.000000	1.379999e+08	6.755589e+08
75%	22.289652	133.000000	2015.000000	1.554341e+08	1.222407e+09
max	32.985763	169.000000	2015.000000	1.839999e+08	2.789712e+09

These are the numerical characteristics of the top ten popular movies. Their runtime seems to be a slightly high compared to other movies. All of them are pretty recent with the exception of Star Wars (not surprisingly). They all are high-budget movies the minimum being around 10 million USD and the median being around 140 million USD. All of them earned more than 70 million USD the highest one reaching almost 2.8 billion USD (Star Wars!).

Popular movie directors Who are the most popular movie directors? (Third question.) There is no unique way to answer this, at least in terms of ranking. I could just choose the directors of most popular movies and rank them according to the movies. But this will be unjust to directors who have many movies which are popular. In the following I have considered 100 most popular movies and counted the directors. Quentin Tarantino and Christopher Nolan come on top, each delivering 5 among the 100 most popular movies.

```
In [24]: dataset.sort_values(by='popularity', ascending=False).head(100).director.value_counts()
```

```
Out[24]: Quentin Tarantino      5
          Christopher Nolan      5
          David Yates            4
          Peter Jackson          4
          Chris Columbus         3
          Francis Lawrence       3
          Gore Verbinski        2
          James Cameron         2
          David Fincher         2
          Jon Favreau           2
          Name: director, dtype: int64
```

1.1.6 2. Revenue

In our discussions on popularity, we saw that popular movies tend to make good money too. In this part we shall explore the distribution of revenue across the movies in the database. But first, let us remember that there are many zero entries in the field. Let us first see how many of them are there.


```
In [25]: dataset.query('revenue_adj == 0').shape
```

```
Out[25]: (5888, 11)
```

There's a lot of zeros! Nevertheless, we still have more than 4000 movies with information on revenue. We create a dataframe for the latter movies and explore.


Revenue proper

```
In [26]: dataset_rev = dataset.query('revenue_adj > 0')
dataset_rev.revenue_adj.hist()
plt.xlabel('Revenue (in 2010 USD)')
plt.ylabel('Counts')
plt.title('Distribution of revenue');
```



We see that here too the distribution is affected by the presence of outliers. To get a visual sense, we use boxplot as before.

```
In [27]: dataset.revenue_adj.plot.box(vert=False, sym='r')
plt.xlabel('Revenue (in 2010 USD)')
plt.title('Boxplot of revenue');
```



So there are many outliers. To get a distribution for the regular entries, I use the parameters of boxplot to separate the outliers. This is the same procedure as was the case with popularity.

```
In [28]: # quartiles
q1_rev = dataset_rev.revenue_adj.describe()['25%']
q3_rev = dataset_rev.revenue_adj.describe()['75%']

# interquartile range
iqr_rev = q3_rev - q1_rev

# fences
fence_low_rev = q1_rev - 1.5 * iqr_rev
fence_high_rev = q3_rev + 1.5 * iqr_rev

# Regular and outlier samples
dataset_rev_reg = dataset_rev.query('revenue_adj > {} and revenue_adj < {}'.format(fence_low_rev, fence_high_rev))
dataset_rev_out = dataset_rev.query('revenue_adj < {} or revenue_adj > {}'.format(fence_low_rev, fence_high_rev))


# shape of the regular dataset
dataset_rev_reg.shape
```

```
Out[28]: (4373, 11)
```

```
In [29]: dataset_rev_reg = dataset_rev.query('revenue_adj > {} and revenue_adj < {}'.format(fence_low_rev, fence_high_rev))
dataset_rev_reg.revenue_adj.plot.hist(bins=20)
```




```
plt.xlabel('Revenue (in 2010 USD)')
plt.title('Distribution of available revenue');
```



We see that unlike popularity histogram, there is no peak in the revenue distribution. It's more skewed. This also answers our fourth question: **How is the revenue distribution?** From the above distribution I wondered if the same is true for the top-grossing movies (outliers).

```
In [30]: dataset_rev_out = dataset_rev.query('revenue_adj < {} or revenue_adj > {}'.format(fence
dataset_rev_out.revenue_adj.plot.hist(bins=20)
plt.xlabel('Revenue (in 2010 USD)')
plt.title('Revenue distribution of top-grossing movies');
```



Yes, the revenue distribution of the top-grossing movies is no different from the rest. But **what are these top-grossing movies?** (Our fifth question.) I give the list of top ten movies in terms of revenue and their characteristics.

Top ten!

In [31]: `dataset.sort_values(by='revenue_adj', ascending=False).head(10)`

```
Out[31]:
```

	popularity	original_title \	cast \
1386	9.432768	Avatar	Sam Worthington Zoe Saldana Sigourney Weaver S...
1329	12.037933	Star Wars	Mark Hamill Harrison Ford Carrie Fisher Peter ...
5231	4.355219	Titanic	Kate Winslet Leonardo DiCaprio Frances Fisher ...
10594	2.010733	The Exorcist	Linda Blair Max von Sydow Ellen Burstyn Jason ...
9806	2.563191	Jaws	
3	11.173104	Star Wars: The Force Awakens	
8889	2.900556	E.T. the Extra-Terrestrial	
8094	1.136610	The Net	
10110	2.631987	One Hundred and One Dalmatians	
4361	7.637767	The Avengers	

9806	Roy Scheider Robert Shaw Richard Dreyfuss Lorr...
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...
8889	Henry Thomas Drew Barrymore Robert MacNaughton...
8094	Sandra Bullock Jeremy Northam Dennis Miller We...
10110	Rod Taylor J. Pat O'Malley Betty Lou Gerson Ma...
4361	Robert Downey Jr. Chris Evans Mark Ruffalo Chr...

	director	runtime \
1386	James Cameron	162
1329	George Lucas	121
5231	James Cameron	194
10594	William Friedkin	122
9806	Steven Spielberg	124
3	J.J. Abrams	136
8889	Steven Spielberg	115
8094	Irwin Winkler	114
10110	Clyde Geronimi Hamilton Luske Wolfgang Reitherman	79
4361	Joss Whedon	143

	genres \
1386	Action Adventure Fantasy Science Fiction
1329	Adventure Action Science Fiction
5231	Drama Romance Thriller
10594	Drama Horror Thriller
9806	Horror Thriller Adventure
3	Action Adventure Science Fiction Fantasy
8889	Science Fiction Adventure Family Fantasy
8094	Crime Drama Mystery Thriller Action
10110	Adventure Animation Comedy Family
4361	Science Fiction Action Adventure

	production_companies	release_date \
1386	Ingenious Film Partners Twentieth Century Fox ...	12/10/09
1329	Lucasfilm Twentieth Century Fox Film Corporation	3/20/77
5231	Paramount Pictures Twentieth Century Fox Film ...	11/18/97
10594	Warner Bros. Hoya Productions	12/26/73
9806	Universal Pictures Zanuck/Brown Productions	6/18/75
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15
8889	Universal Pictures Amblin Entertainment	4/3/82
8094	Columbia Pictures	7/28/95
10110	Walt Disney Productions	1/25/61
4361	Marvel Studios	4/25/12

	release_year	budget_adj	revenue_adj
1386	2009	2.408869e+08	2.827124e+09
1329	1977	3.957559e+07	2.789712e+09
5231	1997	2.716921e+08	2.506406e+09
10594	1973	3.928928e+07	2.167325e+09

9806	1975	2.836275e+07	1.907006e+09
3	2015	1.839999e+08	1.902723e+09
8889	1982	2.372625e+07	1.791694e+09
8094	1995	3.148127e+07	1.583050e+09
10110	1961	2.917944e+07	1.574815e+09
4361	2012	2.089437e+08	1.443191e+09

```
In [32]: dataset.sort_values(by='revenue_adj', ascending=False).head(50).describe()
```

```
Out[32]:
```

	popularity	runtime	release_year	budget_adj	revenue_adj
count	50.000000	50.000000	50.000000	5.000000e+01	5.000000e+01
mean	5.594902	136.560000	1997.700000	1.392524e+08	1.317835e+09
std	4.761581	28.385941	16.044899	8.145784e+07	4.451266e+08
min	0.760503	78.000000	1961.000000	2.372625e+07	9.658933e+08
25%	2.758491	121.250000	1985.500000	7.110709e+07	1.033247e+09
50%	4.950633	136.000000	2002.500000	1.392025e+08	1.145192e+09
75%	6.684794	153.500000	2011.000000	1.886158e+08	1.388085e+09
max	32.985763	201.000000	2015.000000	3.683713e+08	2.827124e+09

Highest grossing movie directors Who directs these movies? (Our sixth question) Following the procedure for popularity we have the following list.

```
In [33]: dataset.sort_values(by='revenue_adj', ascending=False).head(100).director.value_counts()
```


```
Out[33]: Steven Spielberg      7
         Peter Jackson        6
         Michael Bay          4
         George Lucas         4
         David Yates          4
         Christopher Nolan     3
         Sam Raimi             3
         James Cameron         3
         Gore Verbinski        3
         Chris Columbus        3
         Name: director, dtype: int64
```

Definitely not the same as before!

1.1.7 3. Popularity versus Revenue

We saw that the most popular and most revenue generating movies need not overlap. Here we want see if there is any correlation. This is our seventh question: **Is popularity related to revenue?** (It should be positive.)

```
In [34]: dataset_rev.plot(kind='scatter', x='popularity', y='revenue_adj');
         plt.xlabel('Popularity')
         plt.ylabel('Revenue (in 2010 USD)')
         plt.title('Relation between popularity and revenue');
```



As the scatter plot portrays, there is definitely a positive correlation. The outliers have made the correlation a difficult to see. Also, there is a lot of spread. In the following, we ask a related question: How strong is the correlation among the outliers? If a movie is an outlier in popularity, what are the chances it will be an outlier in revenue, and vice versa?

```
In [35]: # sets of popular and revenue outliers
         set_pop_out = set(dataset_pop_out.original_title)
         set_rev_out = set(dataset_rev_out.original_title)
         print('Popular outlier = {} and Revenue outlier = {}'.format(len(set_pop_out), len(set_rev_out)))
```

Popular outlier = 928 and Revenue outlier = 469

```
In [36]: # intersect of sets of popular and revenue outliers
         print('Intersect: {}'.format(len(set_pop_out.intersection(set_rev_out))))
```

Intersect: 330

A popular outlier has less than fifty percent chance (331/930) of being a revenue outlier whereas a revenue outlier has more than fifty percent (331/470) chance of being a popular outlier.

1.1.8 4. Directors

Let's explore the data on directors. (It's a categorical feature compared to the numerical features above.) **Who are the most productive directors?** (Eighth question.)

```
In [37]: dataset.director.value_counts()[:10]
```

```
Out[37]: Woody Allen      45
         Clint Eastwood   34
         Steven Spielberg  29
         Martin Scorsese  28
         Ridley Scott     23
         Ron Howard       22
         Steven Soderbergh 22
         Joel Schumacher  21
         Brian De Palma   20
         Barry Levinson   19
         Name: director, dtype: int64
```

Not surprisingly, Woody Allen (45) comes at top followed by Clint Eastwood (34) and Steven Spielberg (20). Spielberg is closely followed by Martin Scorsese (28). I wondered how many movies does a director generally make?

```
In [38]: dataset.director.value_counts().describe()
```

```
Out[38]: count      5018.000000
         mean        2.138701
         std         2.529060
         min         1.000000
         25%         1.000000
         50%         1.000000
         75%         2.000000
         max         45.000000
         Name: director, dtype: float64
```

It seems that a director typically makes just one movie!

1.1.9 5. Year

**** What's the yearly movie production rate? **** (Question 9.) This is best displayed graphically.


```
In [39]: # select the year column
         year = dataset.release_year

         # value counts
         year_counts = year.value_counts()

         # reverse the value counts series to have the maximum on top
         year_counts_rev = year_counts.reindex(year_counts.index[::-1])

In [40]: year_counts_rev.plot.barh(figsize=(8, 16))
         plt.xlabel('Movies made')
         plt.title('Yearly movie distribution');
```

Yearly movie distribution



So the number movies per year is increasing steadily, though not always monotonically. For example, 2010 had less number of movies made than in year 2009 and 2008. Similarly, 2015 had less movies than 2013 and 2014.

1.1.10 6. Genre

Finally, we address the last question: **Which genres are most popular from year to year?**

```
In [41]: # earliest and latest years of record
earliest = int(dataset.release_year.describe()['min'])
latest = int(dataset.release_year.describe()['max'])

# For each year, collect the genres and find out
# the frequency distribution
pop_gen = {}
for year in np.arange(earliest, latest+1):
    # collect the genre strings of a year as a series
    genres_year = dataset.query('release_year == {}'.format(year))['genres']
    genres = []
    # split and store genre strings in 'genres' list
    for genre in genres_year:
        genres.append(str(genre).split('|'))

    # make a flat list from 'genres' list
    # so that we can determine frequencies
    flat_list = [item for sublist in genres for item in sublist]
    # make the flat list into a pandas series
    # so that we can use the pd.value_counts() function
    genres = pd.Series(flat_list)
    pop_gen[str(year)] = genres.value_counts().idxmax()

# print the most popular genres per year
pop_gen
```

```
Out[41]: {'1960': 'Drama',
          '1961': 'Drama',
          '1962': 'Drama',
          '1963': 'Comedy',
          '1964': 'Drama',
          '1965': 'Drama',
          '1966': 'Drama',
          '1967': 'Comedy',
          '1968': 'Drama',
          '1969': 'Drama',
          '1970': 'Drama',
          '1971': 'Drama',
```




```

'1972': 'Drama',
'1973': 'Drama',
'1974': 'Drama',
'1975': 'Drama',
'1976': 'Drama',
'1977': 'Drama',
'1978': 'Drama',
'1979': 'Drama',
'1980': 'Drama',
'1981': 'Drama',
'1982': 'Drama',
'1983': 'Drama',
'1984': 'Drama',
'1985': 'Comedy',
'1986': 'Drama',
'1987': 'Comedy',
'1988': 'Comedy',
'1989': 'Comedy',
'1990': 'Drama',
'1991': 'Drama',
'1992': 'Drama',
'1993': 'Drama',
'1994': 'Comedy',
'1995': 'Drama',
'1996': 'Drama',
'1997': 'Drama',
'1998': 'Drama',
'1999': 'Drama',
'2000': 'Drama',
'2001': 'Comedy',
'2002': 'Drama',
'2003': 'Comedy',
'2004': 'Drama',
'2005': 'Drama',
'2006': 'Drama',
'2007': 'Drama',
'2008': 'Drama',
'2009': 'Drama',
'2010': 'Drama',
'2011': 'Drama',
'2012': 'Drama',
'2013': 'Drama',
'2014': 'Drama',
'2015': 'Drama'}

```

Drama first! Comedy second. And none else! This is al the more surprising from to the fact that thriller is the most common genre after drama, and not comedy, as we below.

```
In [42]: pd.Series(flat_list).value_counts().plot(kind='pie', figsize=(10,10), label='Distributi
```



Conclusions

In the above analysis I have explored popularity, revenue, their relation, directors, years, and genres in the movie database (TMDB). The prime characteristic of the database is the presence of outlier. In popularity, for example, we saw that a handful of movies reached a score as high as 10 whereas more than 3 out of 4 had a score below 1.

Revenue had a lot of missing values. Of the movies that did have the data, we saw that here too the presence of outliers was prominent. In fact, the distribution seemed exponential, not even having a peak anywhere. In a sense, there is no typical value of revenue.

Next we considered the relation between popularity and revenue with multiple metrics. They are definitely positively correlated on the average, although there is a lot of spread. I compared most popular movies with highest-grossing movies. A popular outlier has less than fifty percent chance (331/930) of being a revenue outlier whereas a revenue outlier has more than fifty percent (331/470) chance of being a popular outlier. I also looked at the directors of these movies. While there is a definite overlap there were some interesting mismatches in each. Quentin Tarantino came out in top as a popular movie maker, but he was not among the top ten revenue generators. Steven

Spielberg came out in top as a revenue generator, but he was not among the top ten popular directors (in my interpretation, of course!).

Just for fun, I also looked at the prolificity of directors. Woody Allen (45) came in top followed by Clint Eastwood (34) and Steven Spielberg (20). Spielberg is closely followed by Martin Scorsese (28).

More and more movies are being made over the years. The progression, however, is not strictly monotonic. For example, 2010 had less number of movies made than in year 2009 and 2008. Similarly, 2015 had less movies than 2013 and 2014.

Let me end the report with a curious observation about genres. I looked at the most popular genre in a year. Drama came first and comedy second. And none else ever made in the list! This is all the more surprising because comedy is not the second most frequent genre over all; it is thriller. Yet, while comedy came out to be most popular genre many times, thriller never showed up. Something to be investigated soon.

1.1.11 Limitations

Any data-science investigation is limited by the quality of the dataset and the analysis. The current investigation is no exception. The dataset for the current investigation has entries for years 1960 to 2015. We could not make any concrete comments about movies made before or afterwards.

The revenue was not listed for almost half of the movies. So the statistics mentioned in relation to revenue, while representative, may not be robust. This is especially true for the outliers for which we may have missed some top-grossing movies.

While the popularity metric made sense in terms of the movies that we know to be popular, its definition was missing in the description of the dataset. While it must be a reasonable function of the various features of a movie, we cannot be sure that it did not favor some types of movies over others.

On the analysis side, I was not able to explore all the features involved in the dataset and their dependency. Given time, it would have been interesting to explore, for example, the influence of cast and directors on popularity, the popularity of different genres, or if a particular genre is getting more and more popular, among others.